

# Linux Kernel Project3 Report

## Memory Management

蒋逸伟 517161910005

## 1 Requirement

Write a module `mtest`, when the module being loaded, it will create a proc file `mtest`. It will accept three parameters as follows:

- `listvma`

It will print the whole virtual memory address of current process. The format is

`start-addr end-addr permission`

- `findpage addr`

Convert the virtual address of current process to physical address. If there is no translation print

`translation not found`

- `writeval addr val`

write a value to a designated address

## 2 Implement

### 2.1 Step1 Implement of the response of write a string to the proc file

```

1 static ssize_t mtest_proc_write(struct file *file ,
2 const char __user * buffer ,
3 size_t count, loff_t * data)
4 {
5     unsigned long val1 , val2;
6     char *tmp = kzalloc((count+1),GFP_KERNEL); //Dynamically alloc memory
7     if(!tmp) return -ENOMEM;
8     if(copy_from_user(tmp, buffer , count)){
9         kfree(tmp);
10        return EFAULT;
11    }
12    if (memcmp(tmp, "listvma", 7) == 0)
13        mtest_list_vma();
14    else if (memcmp(tmp, "findpage", 8) == 0) {
15        if (sscanf(tmp + 8, "%lx", &val1) == 1) {
16            mtest_find_page(val1);
17        }
18    }
19
20    else if (memcmp(tmp, "writeval", 8) == 0) {
21        if (sscanf(tmp + 8, "%lx %lx", &val1, &val2) == 2) {
22            mtest_write_val(val, val2);

```

```

23     }
24 }
25 return count;
26 }

```

## 2.2 Step2 Implement of listvma

The data structure VMA is defined in `include/linux/mm_types.h`. The data member `vm_start` and the `vm_end` record the start address and the end address of the VMA in process space. `vm_flags` describes the VMA permission. The all VMAs of current process is organized by a link list, the `vm_next` and the `vm_prev` records the previous and the next VMA.

```

1 static void mtest_list_vma(void)
2 {
3     struct mm_struct *mm = current->mm;
4     struct vm_area_struct *vma;
5     down_read(&mm->mmap_sem); //Protect the memory
6     for (vma = mm->mmap; vma; vma = vma->vm_next) {
7         printk("start-0x%lx end-0x%lx %c%c%c \n", vma->vm_start, vma->vm_end,
8             vma->vm_flags & VM_READ ? 'r' : '-',
9             vma->vm_flags & VM_WRITE ? 'w' : '-',
10            vma->vm_flags & VM_EXEC ? 'x' : '-');
11     } // Traverse the VMA link_list.
12     up_read(&mm->mmap_sem);
13 }

```

## 2.3 Step3 Implement of findpage

Original x86-64 was limited by 4-level paging to 256 TiB of virtual address space and 64 TiB of physical address space. The 5-level paging is a straight-forward extension of the current page table structure adding one more layer of translation. It bumps the limits to 128 PiB of virtual address space and 4 PiB of physical address space. The 5 level paging type add a table named `p4d` to the origin page tables.

```

1 static void mtest_find_page(unsigned long addr)
2 {
3     p4d_t *p4d; //The p4d is between the pgd and the pud.
4     pud_t *pud;
5     pmd_t *pmd;
6     pgd_t *pgd;
7     pte_t *pte;
8     spinlock_t *ptl;
9     unsigned long kernel_addr;

```

```

10     struct mm_struct *mm = current->mm;
11     down_read(&mm->mmap_sem);
12     struct vm_area_struct *vma;
13     vma = find_vma(mm, addr);
14     struct page *page = NULL;
15     pgd=pgd_offset(mm, addr);
16     if(pgd_none(*pgd) || unlikely(pgd_bad(*pgd))){
17         goto out;
18     }
19     p4d = p4d_offset(pgd, addr);
20     if(p4d_none(*p4d) || unlikely(p4d_bad(*p4d))){
21         goto out;
22     }
23     pud=pud_offset(p4d, addr);
24     if(pud_none(*pud) || unlikely(pud_bad(*pud))){
25         goto out;
26     }
27     pmd = pmd_offset(pud, addr);
28     if(pmd_none(*pmd) || unlikely(pmd_bad(*pmd))){
29         goto out;
30     }
31     pte = pte_offset_map_lock(mm, pmd, addr, &ptl);
32     if(!pte)
33         goto out;
34     if(!pte_present(*pte))
35         goto unlock;
36     page = pfn_to_page(pte_pfn(*pte));
37     if(!page)
38         goto unlock;
39     get_page(page);
40 unlock:
41     pte_unmap_unlock(pte, ptl);
42 out:
43     up_read(&mm->mmap_sem);
44     if (!page)
45         printk("Translation not Found.\n");
46     else
47         {kernel_addr=(unsigned long)page_address(page);
48          kernel_addr+=(addr&~PAGE_MASK);
49          printk("The physical address is 0x%lx\n", kernel_addr);
50         }

```

```
51 }
```

## 2.4 Step 4 Implement of writeeval addr val

To write a value to a virtual address, should satisfy the following points:

- The address from the user space should between the vma->start and the vma->end. The addr+sizeof(val) should less than the vma->end.
- This vma is writable.
- Should translate the virtual address to the physical address.

Now create a translate function:

```
1 static struct page *find_physical_page(struct vm_area_struct *vma, unsigned long addr)
2     pud_t *pud;
3     p4d_t *p4d;
4     pmd_t *pmd;
5     pgd_t *pgd;
6     pte_t *pte;
7     spinlock_t *ptl;
8     struct page *page = NULL;
9     struct mm_struct *mm = vma->vm_mm;
10    pgd=pgd_offset(mm, addr);
11    if(pgd_none(*pgd) || unlikely(pgd_bad(*pgd))){
12        goto out;
13    }
14    p4d = p4d_offset(pgd, addr);
15    if(p4d_none(*p4d) || unlikely(p4d_bad(*p4d))){
16        goto out;
17    }
18    pud=pud_offset(p4d, addr);
19    if(pud_none(*pud) || unlikely(pud_bad(*pud))){
20        goto out;
21    }
22    pmd = pmd_offset(pud, addr);
23    if(pmd_none(*pmd) || unlikely(pmd_bad(*pmd))){
24        goto out;
25    }
26    pte = pte_offset_map_lock(mm, pmd, addr, &ptl);
27    if(!pte)
28        goto out;
29    if(!pte_present(*pte))
```

```

30         goto unlock;
31     page = pfn_to_page(pte_pfn(*pte));
32     if (!page)
33         goto unlock;
34     get_page(page);
35 unlock:
36     pte_unmap_unlock(pte, ptl);
37 out:
38     return page;
39 }

```

Now for the write function:

```

1  static void mtest_write_val(unsigned long addr, unsigned long val)
2  {
3      struct vm_area_struct *vma;
4      struct mm_struct *mm = current->mm;
5      struct page *page;
6      unsigned long kernel_addr;
7      down_read(&mm->mmap_sem);
8      vma = find_vma(mm, addr);
9      if (vma && addr >= vma->vm_start && (addr + sizeof(val)) < vma->vm_end) {
10         if (!(vma->vm_flags & VM_WRITE)) {
11             printk("This vma is not writable for 0x%lx\n", addr);
12             goto out;
13         }
14         page = find_physical_page(vma, addr);
15         if (!page) {
16             printk("page not found for 0x%lx\n", addr);
17             goto out;
18         }
19
20         kernel_addr = (unsigned long)page_address(page);
21         kernel_addr += (addr & ~PAGE_MASK);
22         *(unsigned long *)kernel_addr = val;
23         put_page(page);
24         printk("write 0x%lx to address 0x%lx\n", val, kernel_addr);
25     } else {
26         printk("no vma found for %lx\n", addr);
27     }
28 out:
29     up_read(&mm->mmap_sem);

```

30 }

## 2.5 Step 5 Create a proc file entry

```

1 static struct file_operations proc_mtest_operations = {
2     .write= mtest_proc_write
3 };
4 static struct proc_dir_entry *mtest_proc_entry;
5 static int __init mtest_init(void)
6 {
7     mtest_proc_entry = proc_create("mtest", 0777, NULL,&proc_mtest_operations);
8     if (mtest_proc_entry == NULL) {
9         printk("Error creating proc entry\n");
10        return -1;
11    }
12    printk("create the filename mtest mtest_init sucess\n");
13    return 0;
14 }
15 static void __exit mtest_exit(void)
16 {
17     printk("exit the module.....mtest_exit\n");
18     remove_proc_entry("mtest", NULL);
19 }

```

## 3 Result

### 3.1 listvma

The command is

```
echo "listvma">/proc/mtest
```

```
dmesg
```



```

fish /home/austingulsh/桌面/Linux_kernel_project/3rd_lab
[33127.686014] start-0x7f9b47501000 end-0x7f9b47502000 r--
[33127.686015] start-0x7f9b47502000 end-0x7f9b47503000 rw-
[33127.686015] start-0x7f9b47503000 end-0x7f9b47504000 rw-
[33127.686016] start-0x7ffce2dc2000 end-0x7ffce2df1000 rw-
[33127.686016] start-0x7ffce2dfc000 end-0x7ffce2dff000 r--
[33127.686017] start-0x7ffce2dff000 end-0x7ffce2e00000 r-x
>

```

The result shows the vma of current process.

### 3.2 findpage

The command is

```
echo "findpage0x7f9b47503001">/proc/mtest
dmesg
```

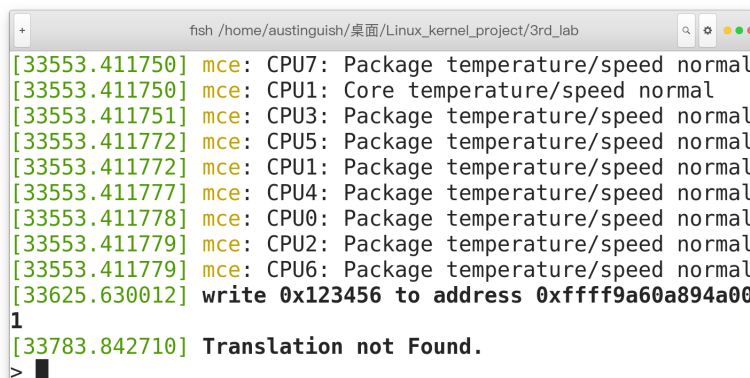


```
fish /home/austinguish/桌面/Linux_kernel_project/3rd_lab

[33127.686015] start-0x7f9b47502000 end-0x7f9b47503000 rw-
[33127.686015] start-0x7f9b47503000 end-0x7f9b47504000 rw-
[33127.686016] start-0x7ffce2dc2000 end-0x7ffce2df1000 rw-
[33127.686016] start-0x7ffce2dfc000 end-0x7ffce2dff000 r--
[33127.686017] start-0x7ffce2dff000 end-0x7ffce2e00000 r-x
[33440.683882] The physical address is 0xffff9a60cdc44001
>
```

If failed (the address is not for current address):

```
echo "findpage0x7f9b475">/proc/mtest
```



```
fish /home/austinguish/桌面/Linux_kernel_project/3rd_lab

[33553.411750] mce: CPU7: Package temperature/speed normal
[33553.411750] mce: CPU1: Core temperature/speed normal
[33553.411751] mce: CPU3: Package temperature/speed normal
[33553.411772] mce: CPU5: Package temperature/speed normal
[33553.411772] mce: CPU1: Package temperature/speed normal
[33553.411777] mce: CPU4: Package temperature/speed normal
[33553.411778] mce: CPU0: Package temperature/speed normal
[33553.411779] mce: CPU2: Package temperature/speed normal
[33553.411779] mce: CPU6: Package temperature/speed normal
[33625.630012] write 0x123456 to address 0xffff9a60a894a00
1
[33783.842710] Translation not Found.
>
```

It shows the physical address of the va 0x7f9b47503001 is 0xffff9a60cdc44001

### 3.3 writeval

The command is

```
echo "writeval0x7f9b474b3001 123456">/proc/mtest
dmesg
```



```

fish /home/austinguish/桌面/Linux_kernel_project/3rd_lab
[33553.411749] mce: CPU5: Core temperature/speed normal
[33553.411750] mce: CPU7: Package temperature/speed normal
[33553.411750] mce: CPU1: Core temperature/speed normal
[33553.411751] mce: CPU3: Package temperature/speed normal
[33553.411772] mce: CPU5: Package temperature/speed normal
[33553.411772] mce: CPU1: Package temperature/speed normal
[33553.411777] mce: CPU4: Package temperature/speed normal
[33553.411778] mce: CPU0: Package temperature/speed normal
[33553.411779] mce: CPU2: Package temperature/speed normal
[33553.411779] mce: CPU6: Package temperature/speed normal
[33625.630012] write 0x123456 to address 0xffff9a60a894a00
1
>

```

If the vma is not writable:

```

fish /home/austinguish/桌面/Linux_kernel_project/3rd_lab
> dmesg | tail
[33553.411751] mce: CPU3: Package temperature/speed normal
[33553.411772] mce: CPU5: Package temperature/speed normal
[33553.411772] mce: CPU1: Package temperature/speed normal
[33553.411777] mce: CPU4: Package temperature/speed normal
[33553.411778] mce: CPU0: Package temperature/speed normal
[33553.411779] mce: CPU2: Package temperature/speed normal
[33553.411779] mce: CPU6: Package temperature/speed normal
[33625.630012] write 0x123456 to address 0xffff9a60a894a00
1
[33783.842710] Translation not Found.
[33940.453039] This vma is not writable for 0x7f9b474d4000
>

```

## 4 Conclusion

- In this lab, I learned the 5-level paging address method, the p4d table is inserted between the pgd and the pud. To convert the address, I should know their mapping relations.
- I read the vma structure carefully and learned the data member of this structure. It records the key info of a process. The vma is an advanced abstract of the physical address. A process can alloc or free a memory more efficiently.
- Learned the spinlock and the semaphore. The macro `pte_offset_map` and the `pte_unmap` are in pairs. It can be used to protect the memory. At `alloc_zeroed_user_highpage`, other thread may update or set the page.